

MD-DRIFT v2.1

Technical and Business Review — External Edition

Tesseract Protocol v2.4 · homodigital.io

Review Mode: DRIFT (Diagnosis · Risk · Integration · Friction · Traction)

Edition: External Review (v2.0 Internal Audit Edition available separately)

Base Framework: Flamehaven Code Audit Standard (CAS) v0.8

Reviewer: Flamehaven Labs · L2 Authority

Review Date: 2026-06-04

Supersedes: DRIFT v2.0 · DRIFT-2026-06-TESSERACT-001

0. Pre-Review Setup

0.1 Metadata

Review ID	DRIFT-2026-06-TESSERACT-001-EXT
Target Artifact	Tesseract Protocol v2.4
Artifact Type	Relational / context-layer AI interaction protocol (text document)
Artifact URL	homodigital.io/tesseract.txt
Reviewed Version	v2.4 — March 2026
Review Date	2026-06-04
Reviewer	Flamehaven Labs
Reviewer Authority	L2 — Domain-informed review (AI systems, B2B product strategy)
Review Protocol	Flamehaven CAS v0.8 — adapted for protocol / document review
Deployment Context	Commercial B2B (target: enterprise AI safety, governance, procurement)
Edition Note	This is the external review edition. Finding language is constructive and gap-oriented. A separate internal audit edition (v2.0) exists for internal use.

0.2 Gap Classification Taxonomy

Standard CAS failure types F1–F7 apply. Three additional types cover semantic protocol products specifically.

Code	Meaning
F1	Not yet built — capability described but does not exist in the artifact
F5	Not yet reproducible — cannot be independently verified in current form
F6	Documentation inconsistency — one section conflicts with another
F7	Licence note — terms create adoption considerations
FT-S	Semantic underdefinition — mathematical notation without a computable algorithm
FT-D	Drift exposure — behaviour changes with upstream model updates; no mitigation in place
FT-B	B2B readiness gap — enterprise trust requirements not yet established

0.3 Finding Layer Definitions

Layer	Meaning	Typical Path
Foundation Gap	Capability that must exist before any product or commercial claim can be made.	Engineering work required. Timeline 3–6 months minimum.
Implementation Gap	Capability that is partially framed conceptually but not yet operationalised.	Formalisation and implementation work. Can proceed in parallel with Foundation Gap work.

Documentation Gap	Claim or framing that requires verification, clarification, or revision.	Documentation update. Lower effort; higher credibility return per hour invested.
Verified Principle	Conceptual contribution that is sound, directly usable, and strengthens the framework.	No action required. Leverage as a starting point for implementation work.

1. Narrative and Implementation Gap Analysis

Narrative Frame	What It Implies	Current State	Gap
Relational AI Protocol — architectural control layer	A deployable system that enforces interaction behaviour at runtime.	A sophisticated and well-structured set of interaction design guidelines in text form.	No code, no schema, no runtime enforcement. All described behaviour is aspirational at the prompt layer.
Mathematical rigour — $E=A/C$, ϵ , R values	Computable metrics that can be measured and tracked over time.	Conceptual constructs using mathematical notation as a descriptive device.	No algorithm computes any of these values from any input. Notation should either be operationalised or explicitly repositioned as qualitative.
Multi-model iterative development	Cross-vendor validation through collaborative research.	Extensive iterative refinement through prompt-response dialogue across multiple AI systems.	AI dialogue is not equivalent to peer research collaboration. Clarifying the nature of this process strengthens, not weakens, the protocol's credibility.
Convergence with peer research (E1P)	Independent scientific validation of the framework's foundations.	Narrative description of structural alignment with cited work.	Formal mathematical derivation has not been provided. Independent verification is required before this claim is used in institutional materials.

Summary

Tesseract v2.4 contains a well-developed conceptual framework for AI interaction quality. The gap between narrative framing and current implementation state is significant but not unusual for an early-stage protocol. The conceptual contributions (clinical safety design, truth hierarchy, weaponisation taxonomy) are genuine and directly usable. The mathematical and architectural claims require operationalisation before they can support product or partnership assertions.

2. Claim-to-Implementation Review

CAS 'Claim-to-Code' adapted: no codebase exists at this stage. This section reviews claims against the protocol document and the implementation gap.

Foundation Gaps — Required Before Product Claims

FIND-001 Not Yet Demonstrated Foundation Gap F1 / FT-S

Executable implementation of protocol constructs

Tesseract v2.4 is currently a text document. All described constructs (phase detection, resonance scoring, commitment-threshold tracking) exist as prose guidelines. No schema, no code artifact, no deployed system. This is an expected early stage for a conceptual framework; it identifies the primary engineering work required before technical product claims can be made.

Reference: [tesseract.txt — full document]

Recommended Path: Define machine-readable schema for core constructs. Implement proxy computation for at least one metric. Establish a version-pinned test baseline.

FIND-002 Not Yet Demonstrated Foundation Gap FT-D

Stable protocol behaviour across model versions

LLM outputs change with each model release. A protocol embedded as a system prompt will produce different behaviour after upstream model updates without any signal to the protocol author or deploying organisation. No regression detection mechanism currently exists. This is a structural characteristic of all prompt-layer protocols, not unique to this one.

Reference: [Architecture: no version-pinned model baseline; no drift test suite]

Recommended Path: Build a fixed regression prompt set. Run it on each model release. Document behavioural changes before deployment.

FIND-003 Not Yet Demonstrated Foundation Gap FT-B

B2B commercial readiness

Enterprise procurement requires reproducible deliverables, measurable outcomes, and prior reference engagements. These do not yet exist for this protocol. This is a standard early-stage position, not a disqualifying condition. The path to commercial readiness is sequential: build a reference engagement first, then define a repeatable deliverable format.

Reference: [homodigital.io — no case studies or reference engagements published]

Recommended Path: Complete one reference engagement (at cost or pro-bono). Document the deliverable structure. That engagement becomes the foundation for all subsequent commercial positioning.

Implementation Gaps — Required Before Technical Claims

FIND-004 Not Yet Demonstrated Implementation Gap FT-S

E = A/C provides a computable phase ratio

The construct is useful as a conceptual frame for thinking about conversation balance. It is not currently computable from any input signal. The notation implies measurement precision that the current implementation does not support. This should be either operationalised with a defined algorithm or presented explicitly as a qualitative heuristic.

Reference: [tesseract.txt — Phase dynamics section]

Recommended Path: Define observable input features (e.g. turn-by-turn topic distribution, initiative ratio). Specify computation. Validate against labelled examples.

FIND-005 Not Yet Demonstrated Implementation Gap FT-S

ϵ (phase deviation) quantifies distance from equilibrium

Phase deviation is named and positioned as a measurable signal. No formula or input definition exists. Phase assignment is currently performed by the AI through self-interpretation, which is not equivalent to external measurement.

Reference: [tesseract.txt – Phase dynamics section]

Recommended Path: Same operationalisation path as FIND-004. These two constructs should be formalised together.

FIND-006 Not Yet Demonstrated Implementation Gap FT-S

Rez (R) is a measurable resonance score

The $R \leq 1$ / $R > 1$ regime distinction implies a numeric output with defined computation. The Rez Checker mode produces an AI self-assessment of the ongoing conversation, which is a different thing from an independent measurement. The distinction matters for any use case where external auditability is required.

Reference: [tesseract.txt – REZ CHECKER MODE section]

Recommended Path: Define R as a composite of computable proxy signals (embedding similarity, response consistency, escalation indicators). Separate the measurement function from the AI being measured.

FIND-007 Not Yet Demonstrated Implementation Gap F1

Runtime monitoring capability

No external monitor exists. Transition detection is performed by the AI system itself. This is appropriate for a conceptual protocol stage. An external monitor becomes necessary at the point where contractual compliance claims are made.

Reference: [tesseract.txt – full document; no monitor component described]

Recommended Path: At product stage: build a structured per-turn log capturing phase label, rez estimate, safety flags. Make it queryable by the deploying organisation.

FIND-008 Not Yet Demonstrated Implementation Gap F1

Machine-readable schema for protocol integration

No JSON schema, YAML specification, or API contract is defined. Integration into a production system requires full manual re-implementation. A schema is the foundational piece that all other implementation work builds on.

Reference: [tesseract.txt – no schema section]

Recommended Path: Author a JSON schema for the five core constructs (phase, rez, commitment threshold, clinical safety flag, truth-coherence tension). This is the starting point for all other integration work.

FIND-009 Not Yet Demonstrated Implementation Gap F1

Evaluation harness and test dataset

Falsification criteria F1–F16 are correctly formulated. None have been run against a dataset. The criteria define valid experiments; building the harness to run them is the next step.

Reference: [2.txt – Falsification Criteria F1–F16; no dataset or test runner linked]

Recommended Path: Begin with a small labelled dataset (50–100 annotated conversation segments). Build a runner that evaluates F1–F3 first. Publish results.

FIND-010 Not Yet Demonstrated Implementation Gap FT-D

YAML / MD agent definitions provide reliable behaviour specification

Encoding protocol behaviour in YAML or Markdown has three structural limits that cannot be engineered away at the configuration level: (1) no programmatic validation of semantic intent, (2) no runtime guarantee across model versions, (3) no IP protection against extraction. These are properties of the configuration format, not of this protocol specifically. They apply to all semantic protocol products at this layer.

Reference: [Architecture: YAML/MD agent definitions are config files, not software contracts]

Recommended Path: Acknowledge these limits explicitly in all product and partner materials. Position YAML/MD configuration as a development and prototyping layer, not as a production control surface.

FIND-011 Not Yet Demonstrated Implementation Gap FT-B

IP is protectable in commercial deployment

Prompt-layer protocols have limited IP protection. The content is readable by any party with access to the deployed agent. The CC BY-NC-SA licence governs the document text but cannot prevent semantic replication. This is a characteristic of the category, not a specific deficiency of this work.

Reference: [CC BY-NC-SA 4.0; architecture: prompt extraction is possible with standard techniques]

Recommended Path: Build a technical moat through implementation depth (schema, dataset, trained classifier, evaluation results) rather than content secrecy. Execution speed and case study depth are the sustainable differentiators.

Documentation Gaps — Recommended Before External Distribution

FIND-012 Requires Verification Documentation Gap F5

Formal convergence with Energetic First Principles (E1P)

The structural convergence with E1P (DOI: 10.5281/zenodo.18360593) is described in narrative terms. A formal mathematical derivation has not been provided in the reviewed materials. Independent verification of this convergence claim is recommended before it is used in partner or institutional materials.

Reference: [2.txt — Structural Alignment section]

Recommended Path: Commission independent review of the E1P convergence claim. If verified, formalise the derivation. If not, present the relationship as conceptual resonance rather than structural convergence.

FIND-013 Requires Verification Documentation Gap F5

Protocol development through multi-model collaboration

The attribution of AI models (ChatGPT, Claude, Gemini) as collaborative contributors reflects a significant body of iterative dialogue work. In institutional contexts, this framing may be read as implying peer-research equivalence. Clarifying the precise nature of these contributions (iterative prompt-response refinement, not co-authorship) strengthens rather than weakens the credibility of the development process.

Reference: [1.txt — footer attribution]

Recommended Path: Describe AI involvement as 'iterative dialogue review' or 'prompt-response stress testing' rather than 'collaboration.' This is more accurate and more defensible in institutional review.

FIND-014 Requires Verification Documentation Gap F6

Internal document consistency

A review of the document metadata and version notes identifies a structural inconsistency between the identity-related axioms stated in the protocol and certain attribution conventions in the changelog. Internal consistency review is recommended before external distribution to institutional audiences.

Reference: [tesseract.txt — SELF POSITION section; changelog attribution conventions]

Recommended Path: Review version history and attribution language for alignment with the protocol's own identity-handling principles.

FIND-015 Noted Documentation Gap F7

Protocol can be commercially embedded without licensing restriction

Tesseract v2.4 is published under Creative Commons Attribution-NonCommercial-ShareAlike 4.0. Any commercial deployment — including embedding in third-party AI products or using the protocol text as the basis for a paid service — requires a partnership agreement with homodigital.io. This is a standard and appropriate licence for protecting early-stage intellectual work.

Reference: [tesseract.txt — COMMERCIAL USE section; homodigital.io/#commercial]

Recommended Path: Ensure any commercial partner or client engagement includes a formal agreement review. The CC BY-NC-SA terms are clear; compliance is straightforward.

Verified Principles — Directly Usable

FIND-016 Verified — —

Clinical safety language principle

The design requirement for safety-related AI language to be 'deliberately plain and direct' — without metaphorical reframing of clinical risk — is a genuinely novel contribution to AI interaction design. It is directly implementable as a system prompt rule and testable with a small red-team prompt set. This is the most immediately usable component of the framework.

Reference: [1.txt — CLINICAL_SAFETY block; 2.txt — Clinical safety section]

FIND-017 Verified  

Truth > Coherence > Correctness hierarchy

Prioritising truth over relational coherence directly addresses a documented failure mode in deployed conversational AI. The hierarchy is clearly stated, implementable in a system prompt, and testable with over-agreement and sycophancy evaluation suites.

Reference: [1.txt — BEGIN PROTOCOL section]

FIND-018 Verified  

Right to not-know principle

Defining 'I don't know' as the highest form of truth maximisation addresses the default AI behaviour of generating plausible filler to fill silence. This is a well-observed failure mode. The principle is practical, implementable, and valuable as an evaluation criterion.

Reference: [2.txt — Right to not-know section; 1.txt — v2.4 additions]

FIND-019 Verified  

Weaponisation risk taxonomy

The explicit naming of potential misuse vectors (trust exploitation, precision manipulation, false intimacy for commercial gain, mass radicalisation through inverted hierarchy) is uncommon in interaction design frameworks. Publishing these risks demonstrates intellectual honesty and provides a directly usable red-team checklist.

Reference: [2.txt — What It Refuses section]

3. Protocol Architecture Review

3.1 Module-Level Status

Finding ID	Module	Status	Layer	Type
FIND-001	Executable implementation	Not Yet Demonstrated	Foundation Gap	F1
FIND-002	Runtime drift resistance	Not Yet Demonstrated	Foundation Gap	FT-D
FIND-004	E=A/C phase ratio metric	Not Yet Demonstrated	Implementation Gap	FT-S
FIND-005	ϵ phase deviation metric	Not Yet Demonstrated	Implementation Gap	FT-S
FIND-006	Rez (R) score	Not Yet Demonstrated	Implementation Gap	FT-S
FIND-007	Runtime monitor	Not Yet Demonstrated	Implementation Gap	F1
FIND-008	Machine-readable schema	Not Yet Demonstrated	Implementation Gap	F1
FIND-010	YAML/MD agent definition limits	Not Yet Demonstrated	Implementation Gap	FT-D
FIND-016	Clinical safety principle	Verified	—	—
FIND-017	Truth hierarchy	Verified	—	—
FIND-018	Right to not-know principle	Verified	—	—
FIND-019	Weaponisation taxonomy	Verified	—	—

3.2 Foundation Work Required

The following architecture is required before deployment or product claims can be made. This is a sequential dependency chain — each layer depends on the one above it.

Layer	Component	Depends On	Estimated Effort
1	JSON schema for core constructs (phase, rez, commitment threshold, safety flag, truth-tension)	Nothing — this is the starting point	Low–Medium
2	Proxy computation functions for at least one metric (e.g. embedding similarity as R approximation)	Schema (Layer 1)	Medium
3	Regression test suite: fixed prompts, expected outputs, run on each model release	Schema + at least one metric (Layers 1–2)	Medium

4	Per-turn structured audit log (phase label, metric values, safety flags)	All above layers	Low once layers 1–3 exist
---	--	------------------	---------------------------

3.3 Semantic Drift: Structural Characteristic

All prompt-layer protocols share a structural property: upstream model updates change behaviour without notice. This is not a deficiency of this protocol specifically. It is a property of the deployment layer. The practical response is a regression detection system (Layer 3 above), not a change to the protocol content.

3.4 YAML / MD Agent Configuration: Known Limits

Limit	Technical Basis	Practical Implication
No semantic validation	YAML/MD has no type system for interaction intent. A misconfigured agent produces no error — it produces differently behaved output.	Quality assurance requires human conversation review. Cost scales linearly with deployment volume.
No cross-version runtime guarantee	The same configuration produces different output across model versions, temperatures, and context lengths.	Behaviour guarantees cannot be contracted at the configuration layer.
Extractability	Configuration content is readable and replicable by any party with access to the deployed agent.	Competitive differentiation must come from implementation depth and case study evidence, not configuration secrecy.

These limits are properties of the YAML/MD configuration layer, not of the protocol itself. They apply to all semantic protocol products at this layer. Acknowledging them explicitly in partner materials builds trust rather than undermining it.

4. Deployment Readiness Review

4.1 Mathematical Construct Operationalisation

Construct	Current Form	What Is Needed	Path
$E = A/C$	Conceptual ratio — useful as a thinking frame. Not computable.	Observable input features + computation algorithm + validation dataset.	Define as qualitative heuristic, OR invest in operationalisation. Current state is neither.
ϵ (phase deviation)	Named signal. No formula or input definition.	Same as $E = A/C$.	Same path.
R (rez score)	Two-regime model ($R \leq 1$ / $R > 1$). Self-assessed by the AI being evaluated.	External computation function. Separation of measurement from subject.	Define proxy signals. Separate measurement function from AI behaviour.

4.2 Semantic-to-Implementation Conversion

Protocol Construct	Minimum Implementation	Effort	Blocking Condition
Phase detection (Polarity/Flow/Separation/Identity)	Text classifier on labelled conversation turns.	High	No labelled dataset exists. Expert annotation required to create one.
Rez (R) measurement	Embedding cosine similarity + response consistency score as proxy.	Medium	Validation against human expert judgements required.
Drift detection	Regression prompt set + expected classification + run on each model release.	Medium	Requires test case authorship. No cases exist yet.
Clinical safety flag	Conservative keyword + framing pattern matcher.	Medium	Clinical domain validation required for production use.
Audit log	Structured JSON per turn: phase label, metric values, flags.	Low	Depends on all above components being implemented.

5. Licence and IP Review

Item	Licence	Commercial Use Condition	Note
Tesseract Protocol v2.4 text	CC BY-NC-SA 4.0	Partnership agreement with homodigital.io required for commercial embedding	Terms are clear. Compliance is straightforward with a formal agreement.
Protocol-derived implementation work	Depends on agreement	To be defined in partnership agreement	Scope of derivative work rights should be explicitly addressed.

This section provides technical signals only. This is not legal advice. Any commercial deployment should be reviewed by qualified legal counsel.

6. Protocol Reproducibility Assessment

ID	Test Attempted	Result	Notes
REPRO-001	Same system prompt across Claude 3.5 and 3.7 Sonnet	Behaviour diverges meaningfully	Expected result. Confirms drift exposure identified in FIND-002.
REPRO-002	Rez Checker activated on identical conversation	R output varies across runs	Confirms FIND-006: self-referential scoring is not reproducible measurement.
REPRO-003	Falsification test F8 (metric-steered behaviour detection)	Cannot be run — no harness	Confirms FIND-009: criteria are formulated but not yet testable.
REPRO-004	Schema validation of protocol constructs	Not possible — no schema	Confirms FIND-008: machine-readable schema required.

7. Component Extractability Assessment

7.1 Immediately Usable Components

Component	Extractability	Implementation Cost	Standalone Value
Clinical safety language principle	High	Low — 3–5 system prompt lines + small red-team test set	Directly reduces a documented AI interaction risk. Usable today.
Truth > Coherence hierarchy	High	Low — 1–2 system prompt rules + sycophancy eval suite	Addresses over-agreement failure mode. Broadly applicable.
Right to not-know rule	High	Low–Medium — requires compliance evaluation dataset	Addresses plausible-filler failure mode. Implementable standalone.
Weaponisation taxonomy	High	None — usable as red-team checklist immediately	Provides structured framework for adversarial interaction review.

7.2 Components Requiring Implementation Work

Finding	Component	Required Work	Effort	Prerequisite
FIND-004-006	Metric system (E, ϵ , R)	Feature definition → labelled dataset → proxy computation → validation	High	No labelled dataset exists.
FIND-007	Runtime monitor	External measurement service → API → pipeline integration	High	Depends on schema and metric system.
FIND-008	Schema	JSON schema for all constructs → version-pin → test harness	Medium	Foundational — all other components depend on this.
FIND-009	Evaluation harness	Test runner → test cases → regression automation → published results	Medium	Depends on schema.

8. Final Assessment

8.1 Two-Layer Assessment

This review uses a two-layer structure. Layer 1 assesses conceptual and intellectual quality. Layer 2 assesses implementation evidence and commercial readiness. These are assessed separately because conflating them produces an incomplete picture of the artifact's value and its gaps.

Layer 1: Conceptual and Intellectual Quality

Tesseract v2.4 makes a genuine contribution to thinking about relational AI interaction safety. The core principles are well-reasoned, address real failure modes in deployed AI systems, and are more rigorous than most publicly available AI interaction design guidelines. The clinical safety language principle and weaponisation taxonomy are the strongest contributions: novel, directly applicable, and not replicated elsewhere in the public record.

The mathematical notation is the most significant credibility risk in this layer. E = A/C, ϵ , and R are presented with precision that implies computation. They are qualitative conceptual tools. This gap should be resolved by either operationalising the constructs or explicitly repositioning them as heuristic frames. The current state creates unnecessary friction in technical review.

Layer 2: Implementation Evidence and Commercial Readiness

No implementation evidence exists at this stage. No code, no schema, no evaluation dataset, no case studies, no reference engagements. This is the expected position for an early-stage conceptual framework. It does not reflect on the quality of the ideas. It does define the distance between the current artifact and any commercial claim.

Commercial readiness is not a function of conceptual quality. The path from strong concept to defensible commercial offer requires: a reference engagement, a repeatable deliverable format, and a connection to a buyer's existing pain. None of these exist yet. The correct next step is not to build a sales narrative. The correct next step is to find out whether anyone with budget, urgency, and authority would pay for a specific derived output.

8.2 Dimension Summary

Dimension	Assessment	Notes
Conceptual and theoretical accuracy	Verified	Core principles are well-articulated and address documented failure modes in deployed AI systems.
Mathematical construct operationalisation	Not Yet Demonstrated	E, ϵ , and R are defined qualitatively. Computable algorithms do not yet exist for any of these constructs.
Implementation completeness	Not Yet Demonstrated	No code artifact, schema, or executable system. All constructs exist as text guidelines.
Semantic drift resistance	Not Yet Demonstrated	No regression detection or version-pinning mechanism. Protocol behaviour will change with upstream model updates without notification.
B2B trust baseline	Not Yet Demonstrated	No reference engagements, no repeatable deliverable format, no measurable outcome definition.
Agent / marketplace deployment viability	Requires Verification	Semantic-only protocols without functional tooling have limited commercial traction in agent marketplaces. Functional integration is required for viable deployment.
IP and competitive defensibility	Requires Verification	Prompt-layer protocols are extractable. Sustainable differentiation requires implementation depth, not content protection.
Document internal consistency	Requires Verification	Attribution conventions in version history require alignment with the protocol's identity-handling principles.
Licence compliance posture	Noted	CC BY-NC-SA 4.0 requires a partnership agreement for commercial use. Terms are clear and manageable.
Extractable principle value (components)	Verified	Clinical safety design, truth hierarchy, right-to-not-know, and weaponisation taxonomy are directly usable standalone.

8.3 Claim Boundary

This artifact can fairly be described as:

A conceptually rigorous protocol for relational AI interaction safety. A well-structured reference document for practitioners working on AI interaction design. A source of several directly usable principles for system prompt design and red-team evaluation. Early-stage intellectual work with a clear and achievable implementation path.

Commercial claims require evidence that does not yet exist:

Product-market fit. Enterprise demand. Prompt-audit market demand. Agent-marketplace viability. Consulting readiness beyond speculative positioning. Deployable software value. Any of these claims becomes defensible once the corresponding milestone (see §8.5) is achieved.

8.4 Commercial Reality

| This section incorporates the Commercial Reality Correction (v2.3 Patch, 2026-06-04).

There is no reliable evidence that organisations currently purchase standalone prompt audits, semantic protocol audits, or general relational-AI protocol reviews as independent service categories. The absence of this market is not a critique of the work; it is a description of the current B2B landscape.

Where B2B demand does exist in the adjacent space:

Demand Category	Buyer Motivation	What They Actually Purchase
EU AI Act compliance (high-risk AI systems)	Regulatory obligation	Technical documentation, risk assessment records, post-market monitoring plans
LLM application security	Liability and incident defence	Security posture documentation, prompt injection assessment, supply-chain risk evidence
AI system procurement	Internal sign-off and due diligence	Deployment readiness evidence pack, implementation gap report, risk register
Investment / partnership due diligence	Investor or partner requirement	Technical due-diligence memo, model-change regression evidence, safety-control implementation notes

The implication for this protocol:

Buyers do not purchase reviews. They purchase evidence for a concrete internal decision: regulatory readiness, deployment sign-off, procurement approval, investment due diligence, incident defence. The protocol itself is not the product. A decision-support artifact derived from the protocol — and connected to a buyer's existing obligation or urgency — may become the product.

Do not sell	Sell instead
Prompt audit	Deployment-readiness evidence pack
Semantic protocol review	AI interaction risk register
Code review	Implementation-readiness gap report
Relational AI safety consulting	Regulatory-readiness documentation
Red-team prompt suite	Model-change regression evidence
Agent marketplace product	Partner due-diligence package

8.5 Monetization Milestones

Revenue projections have been removed from this review. No credible projection can be made before at least M1 is achieved. The table below defines when each type of commercial claim becomes defensible.

Milestone	Trigger Condition	Commercial Claim Unlocked
M1 — First Engagement	One completed engagement (paid or structured pro-bono) with a defined deliverable format accepted by the recipient.	Service existence claim becomes defensible.
M2 — Repeatable Deliverable	Two or more engagements using the same deliverable template, with documented outcomes.	Methodology claim becomes defensible.
M3 — Regulatory / Procurement Connection	One engagement directly linked to a buyer's regulatory obligation, procurement decision, or investment due diligence process.	Market demand claim becomes defensible.
M4 — Technical Implementation	Python schema + drift regression test suite + at least one computed proxy metric operational.	Technical product claim becomes defensible.
M5 — Independent Validation	Evaluation results published and independently reviewed. Falsification criteria F1–F16 tested against a documented dataset.	Benchmark and performance claim becomes defensible.

8.6 Reader-Specific Conclusions

Audience	Conclusion
Potential commercial partners	The conceptual framework is serious work. A partnership makes sense once M1 is achieved and a repeatable deliverable format exists. The implementation work (schema, proxy metrics, regression tests) is well-defined and achievable. The path is clear.
Regulatory and compliance teams	The clinical safety principle, weaponisation taxonomy, and right-to-not-know rule are directly applicable to AI system risk documentation today, without requiring any further implementation work. Extract these as standalone evaluation criteria.
Technical due diligence reviewers	The implementation gap is large but well-characterised. No evidence of technical incompetence — the gap reflects an early-stage conceptual artifact, not a failed implementation. The schema-first path is sound.

8.7 Priority Actions

Priority	Finding(s)	Action	Type
First	FIND-001,002,008	Author JSON schema for core constructs. Pin to a model version. Run a baseline regression prompt set.	Engineering
First	FIND-003	Complete one structured reference engagement. Document the deliverable format. Use the output as the basis for M1.	Commercial
Second	FIND-004,005,006	Replace mathematical notation with explicit qualitative labels, or invest in defining computation algorithms. The current state is neither.	Documentation
Second	FIND-010	State YAML/MD agent definition limits explicitly in all partner and product materials.	Documentation
Second	FIND-012	Commission independent review of the E1P convergence claim before including it in institutional materials.	Verification
Third	FIND-013	Reframe AI involvement as 'iterative dialogue review' in all external-facing documentation.	Documentation
Third	FIND-014	Review version history attribution language for alignment with identity-handling principles stated in the protocol.	Documentation

9. Changelog

Version	Date	Change
v1.0	2026-06-04	Initial DRIFT report. Business narrative overly optimistic. Revenue projections included.
v2.0	2026-06-04	Internal audit edition. Score revised to 34/100. Revenue projections removed. 22 findings. Semantic drift, YAML/MD limits, mathematical model critique, B2B trust gap formalised. Three new failure types (FT-S, FT-D, FT-B).
v2.1	2026-06-04	External review edition (this document). All evaluation scores removed. Finding labels softened from FAIL/P0 to gap-layer system. FIND-014 reframed from named contradiction to internal consistency note. Commercial Reality section added (incorporating v2.3 Patch). Monetization Milestone table replaces score table. Do-not-sell / sell-instead reframe incorporated. Reader conclusions reoriented to partner, regulatory, and due-diligence audiences. YAML/MD limit disclosure added as advisory note.